

Introduction

The sinking of the RMS Titanic in 1912 is a tragic tale of opulence and disaster. On its maiden voyage, the "unsinkable" ship struck an iceberg, leading to the loss of over 1,500 lives. Despite advanced safety features, including watertight compartments, the combination of excessive speed, insufficient lifeboats, and class-based evacuation policies proved fatal. The diverse passenger and crew makeup, ranging from wealthy elites to hopeful immigrants, adds to the poignancy. This maritime catastrophe prompted significant changes in safety regulations and remains a somber reminder of human vulnerability in the face of nature's forces, echoing through history and popular culture.

```
In [1]: # Titanic Survival
```

1 import Necessary Library

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.decomposition import PCA
```

2 import Dataset

```
In [3]: df = pd.read_csv("/kaggle/input/titanic-dataset/Titanic-Dataset.csv")
```

3 Data Analysis

```
In [4]: df.head()
```

Out[4]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.25
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.28
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.92
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.10
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.05

In [5]:

```
df.tail()
```

Out[5]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	13.00
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	30.00
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607	23.45
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	30.00
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	7.73

In [6]:

```
df.shape
```

Out[6]: (891, 12)

In [7]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column        Non-Null Count  Dtype  
---  -
 0   PassengerId   891 non-null    int64  
 1   Survived      891 non-null    int64  
 2   Pclass        891 non-null    int64  
 3   Name          891 non-null    object  
 4   Sex           891 non-null    object  
 5   Age           714 non-null    float64 
 6   SibSp         891 non-null    int64  
 7   Parch         891 non-null    int64  
 8   Ticket        891 non-null    object  
 9   Fare          891 non-null    float64 
10   Cabin         204 non-null    object  
11   Embarked      889 non-null    object  
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

In [8]: `df.dtypes`

```
Out[8]: PassengerId    int64
Survived             int64
Pclass               int64
Name                 object
Sex                  object
Age                  float64
SibSp                int64
Parch                int64
Ticket               object
Fare                 float64
Cabin                object
Embarked             object
dtype: object
```

In [9]: `df.describe()`

```
Out[9]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.200000
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910000
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.450000
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329167

In [10]: `df.ndim`

Out[10]: 2

4 Data cleaning and Preprocessing:

```
In [11]: # Data preprocessing
df.drop(['Name', 'Ticket', 'Fare', 'Cabin', 'PassengerId'], axis=1, inplace=True)
df['Family'] = df['SibSp'] + df['Parch'] + 1
df.drop(['SibSp', 'Parch'], axis=1, inplace=True)
df['Sex'].replace(['male', 'female'], [0, 1], inplace=True)
df.Embarked.replace(['S', 'C', 'Q'], [1, 2, 3], inplace=True)
```

```
In [12]: df.isnull().sum()
```

```
Out[12]: Survived      0
Pclass      0
Sex          0
Age        177
Embarked     2
Family       0
dtype: int64
```

```
In [13]: df['Age'].fillna(df['Age'].median(), inplace=True)
df['Embarked'].fillna(1, inplace=True)
```

```
In [14]: df.columns
```

```
Out[14]: Index(['Survived', 'Pclass', 'Sex', 'Age', 'Embarked', 'Family'], dtype='object')
```

```
In [15]: df.head()
```

```
Out[15]:
```

	Survived	Pclass	Sex	Age	Embarked	Family
0	0	3	0	22.0	1.0	2
1	1	1	1	38.0	2.0	2
2	1	3	1	26.0	1.0	1
3	1	1	1	35.0	1.0	2
4	0	3	0	35.0	1.0	1

```
In [16]: # Define features and target variable
x = df[['Pclass', 'Sex', 'Age', 'Embarked', 'Family']]
y = df['Survived']
```

```
In [17]: df.isnull().any()
```

```
Out[17]: Survived      False
Pclass      False
Sex          False
Age          False
Embarked     False
Family       False
dtype: bool
```

```
In [18]:
```

```
df.dtypes
```

```
Out[18]: Survived      int64  
Pclass      int64  
Sex         int64  
Age         float64  
Embarked    float64  
Family      int64  
dtype: object
```

```
In [19]: df['Survived'].value_counts()
```


```
Out[19]: Survived  
0      549  
1      342  
Name: count, dtype: int64
```

5| Data visualisation

EDA (Exploratory Data Analysis)


5.1 displot

```
In [20]: sns.displot(df, kde=False, bins=30)  
plt.show()
```




5.2 Scatter Plot

```
In [21]: df.plot(kind='scatter', x='Age', y='Embarked')
plt.show()
```



```
In [22]: sns.set_style("whitegrid");
sns.FacetGrid(df, hue="Survived").map(plt.scatter, "Age", "Embarked").add_legend()
plt.show();
```

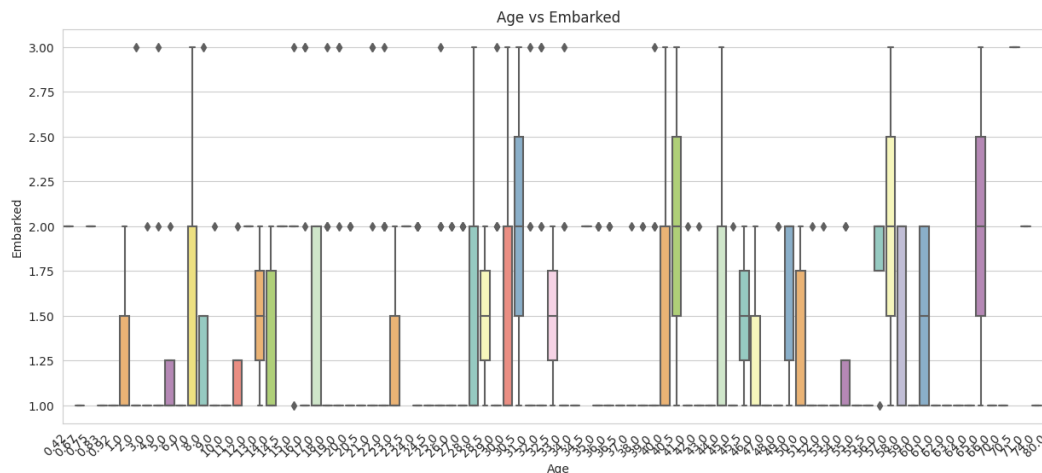


5.3 boxplot

```
In [23]: # sepal_length vs sepal_width boxplot

plt.figure(figsize=(15, 6))
sns.boxplot(x='Age', y='Embarked', data=df, palette='Set3')
plt.title('Age vs Embarked')
plt.xlabel('Age')
```

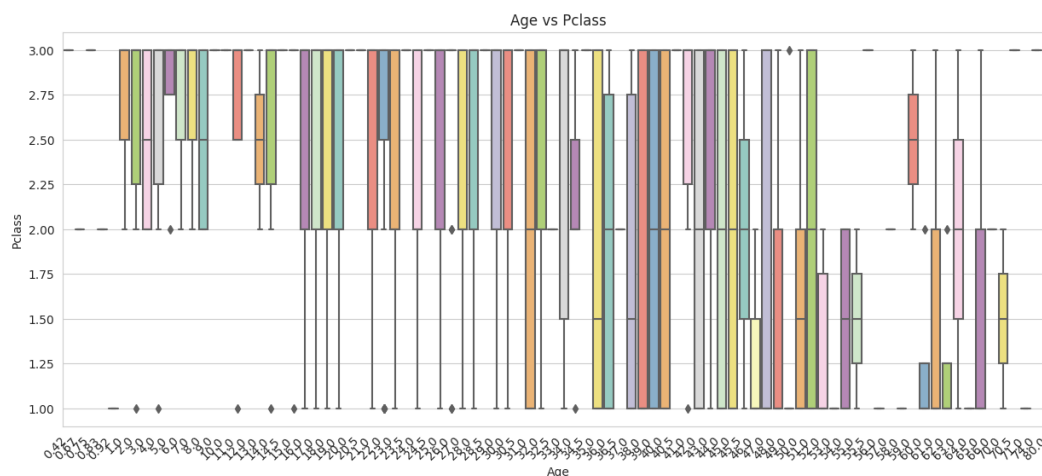
```
plt.ylabel('Embarked')
plt.xticks(rotation=45, ha='right')
plt.show()
```



In [24]:

```
# petal_length vs petal_width boxplot

plt.figure(figsize=(15, 6))
sns.boxplot(x='Age', y='Pclass', data=df, palette='Set3')
plt.title('Age vs Pclass')
plt.xlabel('Age')
plt.ylabel('Pclass')
plt.xticks(rotation=45, ha='right')
plt.show()
```




5.4 countplot

In [25]:

```
# sepal_length


plt.figure(figsize=(15, 8))
sns.countplot(x='Age', data=df, palette='muted')
plt.title('Age')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.xticks(rotation=45, ha='right')
plt.show()
```






In [26]:


```
# sepal_width
plt.figure(figsize=(15, 8))
sns.countplot(x='Pclass', data=df, palette='muted')
plt.title('Pclass')
plt.xlabel('Pclass')
plt.ylabel('Frequency')
plt.xticks(rotation=45, ha='right')
plt.show()
```



In [27]:


```
# petal_length
plt.figure(figsize=(15, 8))
sns.countplot(x='Sex', data=df, palette='muted')
plt.title('Sex')
plt.xlabel('Sex')
plt.ylabel('Frequency')
plt.xticks(rotation=45, ha='right')
plt.show()
```





In [28]:

```
# sepal_width
plt.figure(figsize=(15, 8))
sns.countplot(x='Embarked', data=df, palette='muted')
plt.title('Embarked')
plt.xlabel('Embarked')
plt.ylabel('Frequency')
plt.xticks(rotation=45, ha='right')
plt.show()
```




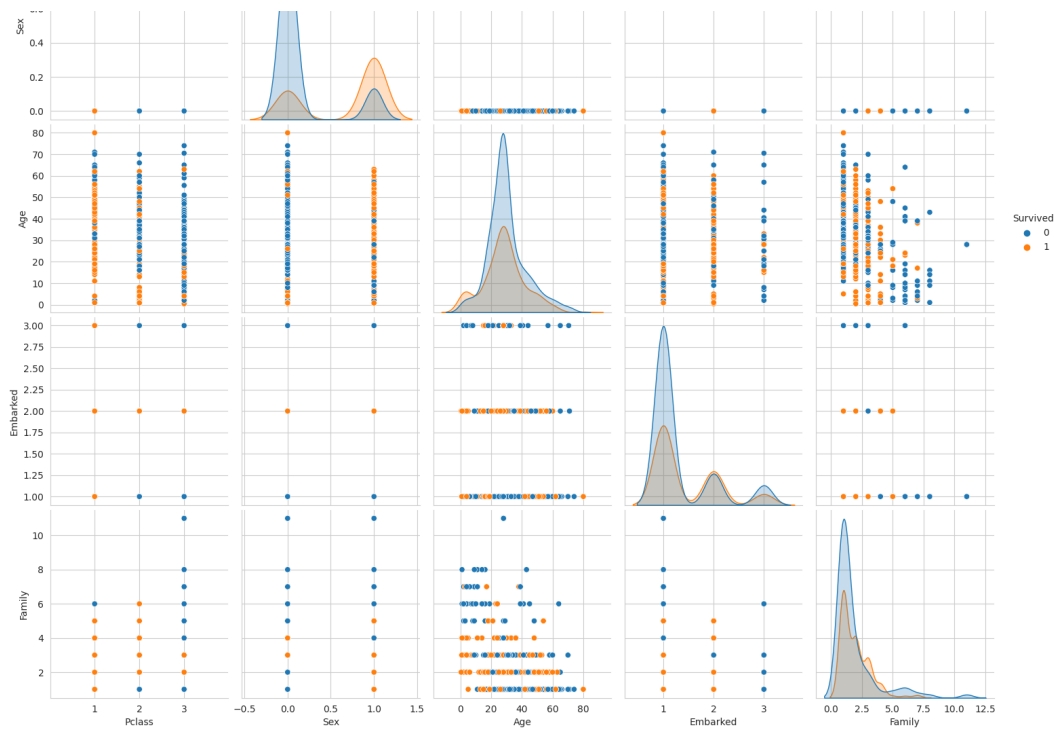
5.5 pairplot

In [29]:

```
sns.set_style("whitegrid")
sns.pairplot(df, hue="Survived", size=3)
plt.show()
```

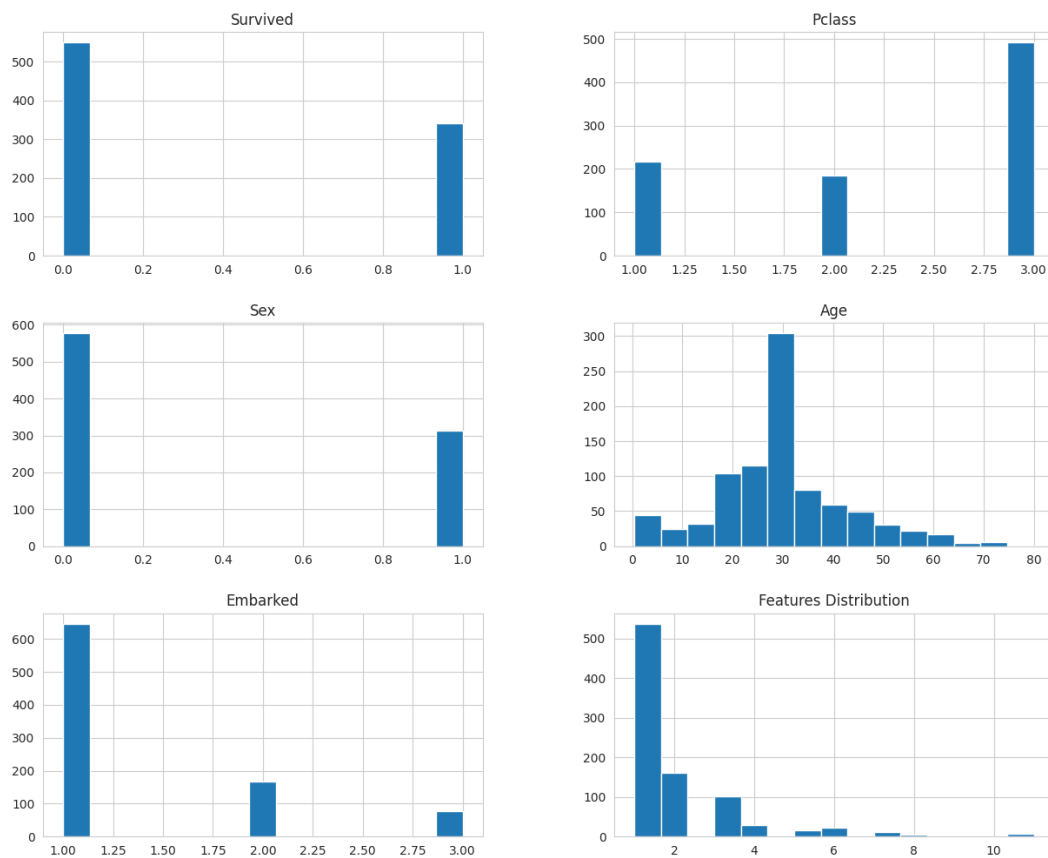
/opt/conda/lib/python3.10/site-packages/seaborn/axisgrid.py:2095: UserWarning: The `size` parameter has been renamed to `height`; please update your code.
warnings.warn(msg, UserWarning)





5.6 Hist Plot


```
In [30]: df.hist(figsize=(15,12),bins = 15)
plt.title("Features Distribution")
plt.show()
```



5.7 violinplot


```
In [31]: sns.violinplot(x="Survived",y="Pclass", data=df, size = 8)
```

```
Out[31]: <Axes: xlabel='Survived', ylabel='Pclass'>
```




```
In [32]: sns.violinplot(x="Survived",y="Sex", data=df, size = 8)
```

```
Out[32]: <Axes: xlabel='Survived', ylabel='Sex'>
```




```
In [33]: sns.violinplot(x="Survived",y="Age", data=df, size = 8)
```

Out[33]: <Axes: xlabel='Survived', ylabel='Age'>




In [34]: `sns.violinplot(x="Survived",y="Embarked", data=df, size = 8)`

Out[34]: <Axes: xlabel='Survived', ylabel='Embarked'>



In [35]: `sns.violinplot(x="Survived",y="Family", data=df, size = 8)`

Out[35]: <Axes: xlabel='Survived', ylabel='Family'>



5.8 Pie Plot

```
In [36]: # ax=plt.subplots(1,1,figsize=(10,8))
# df['species'].value_counts().plot.pie(explode=[0.1,0.1,0.1],autopct='%1.1f%%'
# plt.title("Iris Species %")
# plt.show()
```

6 | Split the Dataset

```
In [37]: from sklearn.model_selection import train_test_split
```

```
In [38]: df.head()
```

```
Out[38]:
```

	Survived	Pclass	Sex	Age	Embarked	Family
0	0	3	0	22.0	1.0	2
1	1	1	1	38.0	2.0	2
2	1	3	1	26.0	1.0	1
3	1	1	1	35.0	1.0	2
4	0	3	0	35.0	1.0	1

```
In [39]: X = df[["Pclass", "Sex", "Age", "Embarked", "Family"]]
```

```
In [40]: y = df['Survived']
```

```
In [41]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.2, random

In [42]: X_train.shape, X_test.shape, y_train.shape, y_test.shape

Out[42]: ((712, 5), (179, 5), (712,), (179,))
```

7 | PCA (Principal Component Analysis)

```
In [43]: from sklearn.decomposition import PCA
```

```
In [44]: pca = PCA(n_components=2)
```

```
In [45]: pca
```

```
Out[45]: PCA(n_components=2)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [46]: X_pca = pca.fit_transform(X)
```

```
In [47]: X_pca[0]
```

```
Out[47]: array([-7.37313857, -0.14876855])
```

```
In [48]: print("Explained Variance Ratio:")
print(pca.explained_variance_ratio_)
```

```
Explained Variance Ratio:
[0.97870884 0.01416231]
```

```
In [49]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X = scaler.fit_transform(X)
```

```
In [50]: X[0]
```

```
Out[50]: array([1.          , 0.          , 0.27117366, 0.          , 0.1          ])
```

```
In [51]: X[1]
```

```
Out[51]: array([0.          , 1.          , 0.4722292, 0.5          , 0.1          ])
```

(1) KNN

```
In [52]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
```

```
In [53]: knn_classifier = KNeighborsClassifier(n_neighbors=3)
```

```
In [54]: knn_classifier.fit(X_train, y_train)
```

```
Out[54]: KNeighborsClassifier(n_neighbors=3)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [55]: train_predictions = knn_classifier.predict(X_train)

train_accuracy1 = accuracy_score(y_train, train_predictions)
```

```
In [56]: test_predictions = knn_classifier.predict(X_test)

test_accuracy1 = accuracy_score(y_test, test_predictions)
```

```
In [57]: print(f"Training Accuracy: {train_accuracy1}")
print(f"Testing Accuracy: {test_accuracy1}")
```

Training Accuracy: 0.8553370786516854

Testing Accuracy: 0.770949720670391

(2) Naive Bayes classifier

```
In [58]: from sklearn.naive_bayes import GaussianNB
from sklearn.naive_bayes import BernoulliNB
from sklearn.naive_bayes import MultinomialNB

from sklearn import metrics
```

```
In [59]: # GaussianNB
```

```
In [60]: G_classifier = GaussianNB()
```

```
In [61]: G_classifier.fit(X_train, y_train)
```

```
Out[61]: GaussianNB()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [62]: train_predictions = G_classifier.predict(X_train)

train_accuracy21 = accuracy_score(y_train, train_predictions)
```

```
In [63]: test_predictions = G_classifier.predict(X_test)

test_accuracy21 = accuracy_score(y_test, test_predictions)
```

```
In [64]: print(f"Training Accuracy: {train_accuracy21}")
print(f"Testing Accuracy: {test_accuracy21}")
```

Training Accuracy: 0.7935393258426966
Testing Accuracy: 0.7653631284916201

```
In [65]: # BernoulliNB
```

```
In [66]: B_classifier = BernoulliNB()
```

```
In [67]: B_classifier.fit(X_train, y_train)
```

Out[67]: BernoulliNB()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [68]: train_predictions = B_classifier.predict(X_train)

train_accuracy22 = accuracy_score(y_train, train_predictions)
```

```
In [69]: test_predictions = G_classifier.predict(X_test)

test_accuracy22 = accuracy_score(y_test, test_predictions)
```

```
In [70]: print(f"Training Accuracy: {train_accuracy22}")
print(f"Testing Accuracy: {test_accuracy22}")
```

Training Accuracy: 0.7879213483146067
Testing Accuracy: 0.7653631284916201

```
In [71]: # MultinomialNB
```

```
In [72]: M_classifier = MultinomialNB()
```

```
In [73]: M_classifier.fit(X_train, y_train)
```

Out[73]: MultinomialNB()

In a Jupyter environment, please rerun this cell to show the HTML representation

or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [74]: train_predictions = M_classifier.predict(X_train)

train_accuracy23 = accuracy_score(y_train, train_predictions)
```

```
In [75]: test_predictions = M_classifier.predict(X_test)

test_accuracy23 = accuracy_score(y_test, test_predictions)
```

```
In [76]: print(f"Training Accuracy: {train_accuracy23}")
print(f"Testing Accuracy: {test_accuracy23}")
```

Training Accuracy: 0.7865168539325843
Testing Accuracy: 0.7821229050279329

```
In [77]: # GaussianNB
# BernoulliNB
# MultinomialNB

# Being the best of them | GaussianNB |
```

(3) Decision Tree

```
In [78]: from sklearn.tree import DecisionTreeClassifier
```

```
In [79]: clf = DecisionTreeClassifier()
```

```
In [80]: clf.fit(X_train, y_train)
```

Out[80]: DecisionTreeClassifier()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [81]: train_predictions = clf.predict(X_train)

train_accuracy3 = accuracy_score(y_train, train_predictions)
```

```
In [82]: test_predictions = clf.predict(X_test)

test_accuracy3 = accuracy_score(y_test, test_predictions)
```

```
In [83]: print(f"Training Accuracy: {train_accuracy3}")
print(f"Testing Accuracy: {test_accuracy3}")
```

Training Accuracy: 0.9339887640449438

Testing Accuracy: 0.7932960893854749

(4) Random Forest

```
In [84]: from sklearn.ensemble import RandomForestClassifier
```

```
In [85]: rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
In [86]: rf_classifier.fit(X_train, y_train)
```

```
Out[86]: RandomForestClassifier(random_state=42)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [87]: train_predictions = rf_classifier.predict(X_train)
         train_accuracy4 = accuracy_score(y_train, train_predictions)
```

```
In [88]: test_predictions = rf_classifier.predict(X_test)
         test_accuracy4 = accuracy_score(y_test, test_predictions)
```

```
In [89]: print(f"Training Accuracy: {train_accuracy4}")
         print(f"Testing Accuracy: {test_accuracy4}")
```

Training Accuracy: 0.9339887640449438

Testing Accuracy: 0.8212290502793296

(5) Boosting Algorithm

```
In [90]: from sklearn.ensemble import AdaBoostClassifier
```

```
In [91]: base_classifier = DecisionTreeClassifier(max_depth=3)
```

```
In [92]: adaboost_classifier = AdaBoostClassifier(base_classifier, n_estimators=50, random_state=42)
```

```
In [93]: adaboost_classifier.fit(X_train, y_train)
```

```
Out[93]: AdaBoostClassifier(estimator=DecisionTreeClassifier(max_depth=3),
                             random_state=42)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [94]: train_predictions = adaboost_classifier.predict(X_train)

train_accuracy5 = accuracy_score(y_train, train_predictions)
```

```
In [95]: test_predictions = adaboost_classifier.predict(X_test)

test_accuracy5 = accuracy_score(y_test, test_predictions)
```

```
In [96]: print(f"Training Accuracy: {train_accuracy5}")
print(f"Testing Accuracy: {test_accuracy5}")
```

Training Accuracy: 0.925561797752809

Testing Accuracy: 0.7877094972067039

(6).SVM

```
In [97]: from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
```

```
In [98]: scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [99]: svm_classifier = SVC(kernel='linear', C=1.0)
```

```
In [100]: svm_classifier.fit(X_train, y_train)
```

Out[100]: SVC(kernel='linear')

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [101]: train_predictions = svm_classifier.predict(X_train)

train_accuracy6 = accuracy_score(y_train, train_predictions)
```

```
In [102]: test_predictions = svm_classifier.predict(X_test)

test_accuracy6 = accuracy_score(y_test, test_predictions)
```

```
In [103]: print(f"Training Accuracy: {train_accuracy6}")
print(f"Testing Accuracy: {test_accuracy6}")
```

Training Accuracy: 0.7879213483146067

Testing Accuracy: 0.7821229050279329

(7). Logistic Regression

```
In [104]:
```

```
from sklearn import linear_model
```

```
In [105... lrg = linear_model.LogisticRegression()
```

```
In [106... lrg.fit(X_train, y_train)
```

```
Out[106... LogisticRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [107... train_predictions = lrg.predict(X_train)

train_accuracy7 = accuracy_score(y_train, train_predictions)
```

```
In [108... test_predictions = lrg.predict(X_test)

test_accuracy7 = accuracy_score(y_test, test_predictions)
```

```
In [109... print(f"Training Accuracy: {train_accuracy7}")
print(f"Testing Accuracy: {test_accuracy7}")
```

Training Accuracy: 0.8033707865168539

Testing Accuracy: 0.8044692737430168

(8).Linear Regression

```
In [110... from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
```

```
In [111... model = LinearRegression()
```

```
In [112... model.fit(X_train, y_train)
```

```
Out[112... LinearRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [113... train_predictions = clf.predict(X_train)

train_accuracy8 = accuracy_score(y_train, train_predictions)
```

/opt/conda/lib/python3.10/site-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but DecisionTreeClassifier was fitted with feature names
warnings.warn(

```
In [114...
```

```
test_predictions = clf.predict(X_test)

test_accuracy8 = accuracy_score(y_test, test_predictions)
```

/opt/conda/lib/python3.10/site-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but DecisionTreeClassifier was fitted with feature names
warnings.warn(

```
In [115... print(f"Training Accuracy: {train_accuracy8}")
print(f"Testing Accuracy: {test_accuracy8}")
```

Training Accuracy: 0.3890449438202247
Testing Accuracy: 0.4134078212290503

(9).Gradient Boosting Machines (GBM)

```
In [116... from sklearn.ensemble import GradientBoostingClassifier
```

```
In [117... model = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, max_de
```

```
In [118... model.fit(X_train, y_train)
```

Out[118... GradientBoostingClassifier(random_state=42)
In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [119... train_predictions = model.predict(X_train)

train_accuracy9 = accuracy_score(y_train, train_predictions)
```

```
In [120... test_predictions = model.predict(X_test)

test_accuracy9 = accuracy_score(y_test, test_predictions)
```

```
In [121... print(f"Training Accuracy: {train_accuracy9}")
print(f"Testing Accuracy: {test_accuracy9}")
```

Training Accuracy: 0.8707865168539326
Testing Accuracy: 0.8156424581005587

Random Forest, Decision Tree, Gradient Boosting Machines (GBM), Algorithm is the best accuracy

(GradientBoostingClassifier)

accuracy mean = 0.90

accuracy mean 0.90



THANK YOU!


10 | Hierarchical Clustering


```
In [122... from scipy.cluster.hierarchy import linkage, dendrogram, fcluster
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
```

```
In [123... scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
In [124... linkage_matrix = linkage(X_scaled, method='ward')
```

```
In [125... plt.figure(figsize=(12, 6))
dendrogram(linkage_matrix, labels=df['Survived'].values, orientation='top', di
plt.title('Titanic Survival Dendrogram')
plt.xlabel('Survived')
plt.ylabel('Distance')
plt.show()
```





In []: